

Diffusing Computation

Using Spanning Tree Construction for Solving Leader Election

- Root is the leader
- In the presence of faults,
 - There may be multiple trees
 - Multiple leaders
- After recovery
 - There is a unique tree
 - There is a unique leader
- A spanning tree maintenance algorithm is a nonmasking leader election algorithm
 - Make it masking fault-tolerant

Using Spanning Tree Construction for Solving Leader Election

- Spanning tree provides a nonmasking fault-tolerant solution
 - Eventually, there is one leader
 - During recovery,
 - The safety specification could be violated because there could be multiple leaders

Designing Masking Fault-Tolerant Solution for Leader Election

- Need to ensure that there is at most one leader during recovery
- Can be achieved if a node makes sure that before it declares itself to be the leader, there is no other leader
 - If this node were to become the true leader, the tree would be reconstructed to be rooted at that node
 - Check if the tree has been reconstructed
 - I.e., check if all nodes are in the tree rooted at the current node

How?

- When j changes $P.j$ to j
 - Check if all nodes have set their root value to j ?
 - Achieved through diffusing computation
 - Similar to one of the termination detection algorithms (next topic)

Diffusing Computation

- Initiation
 - The node that initiates the diffusing computation sends the diffusing computation message to all its children
- Propagation
 - A node propagates the diffusing computation when it receives it (for the first time)
 - To propagate, the node sends the diffusion message to all its children
 - A node can check some predicate (condition) during propagation

Diffusing Computation

- Completion
 - A node completes the diffusing computation iff
 - All its children complete the diffusing computation
 - All its neighbors have received (propagated) that diffusing computation
 - A node can also check some predicate during completion
 - The results from this condition are propagated towards the root;
- The diffusing computation completes when the root completes the diffusing computation
 - If the condition checked by the root evaluates to true, we say that the diffusing computation completes successfully
 - Else, we say that the diffusing computation fails/completes unsuccessfully

Using Spanning Tree for Leader Election

- When a node is about to become leader
 - It starts a diffusing Computation
 - If it completes successfully, declare itself to be the leader
- Goal of the diffusing computation is to check if the tree is formed at the initiator node
 - In other words, check if root value of all nodes is equal to that of the initiator

Property

- If the tree is not formed then
 - There is at least one node j such that
 - j is not in the tree
 - But j has a neighbor that is in the tree

Issues

- Issues
 - Multiple nodes may start diffusing computation
 - Node with higher ID should win (I.e., its diffusing computation should complete successfully)
 - Node with lower ID should lose.
 - Failures during diffusing computation
 - Fail the current diffusing computation

Upon completion

- Upon completing an unsuccessful diffusing computation
 - Start another one if the node is still likely to be the leader ($P.j = j$)
 - The diffusing computation may have failed due to faults
- Use sequence number to distinguish between different diffusing computations initiated by the same node

Actions

- Init

P.j = j



- Phase.j = prop
- Sn.j = newseq()
- Res.j = true
 - (Result = result of the diffusing computation)
 - Currently set to true, to be read only after the diffusing computation completes

Actions

- Propagation

$\text{Root.j} = \text{root.}(P.j) \wedge \text{sn.j} \neq \text{sn.}(P.j)$

→

- If ($\text{phase.}(P.j) = \text{prop}$)
 - $\text{Phase.j} = \text{prop}, \text{res.j} = \text{true}$
- Else
 - $\text{Res.j} = \text{false}$
 - » // Fail this diffusing computation

Actions

- Completion

Phase.j = prop \wedge

$\forall k : k \in \text{Neighbors.j} : \text{root.j} = \text{root.k} \wedge \text{sn.j} = \text{sn.k} \wedge$

$\forall k : k \in \text{Ch.j} : \text{phase.k} = \text{comp}$

\rightarrow

$\text{res.j} = \forall k : k \in \text{Neighbors.j} \cup \{j\} : \text{res.k}$

Phase.j = comp

If (P.j = j \wedge \neg res.j) Init(j) // Initiate a new diffusing computation

Actions

- Aborting Diffusing Computation

Phase.j = comp, res.j = false

If $(P.j \in \text{Neighbors.j})$ res.(P.j) = false

// Last part when j changes its parent. With respect to `old' parent.

Combining with Tree Algorithm

- When j executes tree correction action 1 (red color propagation)
 - Abort(j)
- When j executes tree correction action 2 (changing color to green)
 - Init(j)
- When j changes tree
 - Abort(j)

Application in Mutual Exclusion

- Allow only tree root to generate the new token after faults
 - When a node becomes a root, set $h.j = j$ thereby permitting it to possibly generate a token
 - Don't send this token to other processes unless you ensure that you can actually generate the token
- If token is to be generated (safely) at the root then
 - All nodes must be in the same tree // done already with current task
 - For all nodes: $h.j = P.j$ // condition to be checked at each node
- Init/Prop actions same
- Complete action changed as:

Actions

- Completion

Phase.j = prop \wedge

$\forall k : k \in \text{Neighbors.j} : \text{root.j} = \text{root.k} \wedge \text{sn.j} = \text{sn.k} \wedge$

$\forall k : k \in \text{Ch.j} : \text{phase.j} = \text{comp}$

\rightarrow

$\text{res.j} = \forall k : k \in \text{Neighbors.j} \cup \{j\} : (\text{res.k} \ \& \ \text{h.k} = \text{P.k})$

Phase.j = comp

If ($\text{P.j} = j \wedge \neg \text{res.j}$) Init(j) // Initiate a new diffusing computation

Diffusing Computation

- Can be performed in the absence of a tree
 - Tree is formed during diffusing computation
 - The initiator sends the first diffusion message
 - If node, say j , receives a diffusion from k then
 - If this is NOT the first diffusion message
 - » Send a reply to k
 - If this is the first diffusion message then
 - » Send the diffusion to all neighbors
 - » Send a reply to k after replies from ALL neighbors is received
 - » A condition can be checked based on the replies received

Diffusing Computation (continued)

- The node from which the diffusion message is received for the first time is the parent of that node.
 - Several problems can be solved by taking appropriate actions during propagation and completion of diffusion
 - Detecting global state
 - Termination detection

Applications of Diffusing Computations and other Terminology

- Route Discovery of protocols such as DSR
- Viral computations
- Global snapshots

Scalpel vs Hammer